Private Benchmark Verification — Codebook v3.1

General coding rules

Unit & cutoff	Code one benchmark version at the frozen corpus cutoff; tie evidence to that version.
Missing evidence	Use “Not reported” or “None documented” when supporting evidence is not located. Do not infer absence.
Multi-select fields	Select every documented option. “None documented” is exclusive of positive options.
Conditional fields	Use N/A only when the question is logically inapplicable, not when evidence is missing.
Analysis role	Descriptive = architecture/context (useful for stratification). Comparison = pre-specified cross-benchmark variable. Neither implies a score; do not sum codes.
Verification vs. verifiability	An available way to check a claim is not the same as an external check that actually occurred.
Traceability	Save the source URL and exact supporting passage for each variable. Add a note only for Other, Unclear, or borderline cases.

Response type	Analysis role	Variable	Coding question	Allowed codes		Decision rule	Code definitions
		1. Privacy architecture & durability					
Multi-select	Descriptive	**Privacy rationale**	Why does the benchmark documentation say information is withheld?	Contamination/leakage; adaptive gaming/overfitting; confidentiality/legal; proprietary/IP; fair or blind evaluation; security/safety; not reported		Code only stated rationales; do not infer motivation.	Contamination/leakage — prevent test content from entering training/retrieval data; adaptive gaming/overfitting — prevent repeated probing or benchmark-specific tuning; confidentiality/legal — restricted because of privacy, licensing, regulation, or legal duty; proprietary/IP — protected as non-public commercial or intellectual property; fair or blind evaluation — keep participants from seeing material needed for an unbiased comparison; security/safety — disclosure could create security or misuse risk; other — stated rationale not covered above; not reported — no rationale located.
Multi-select	Descriptive	**Protected component**	Which consequential parts of the evaluation are intentionally withheld?	Items/tasks; labels/reference answers; rubric/judge/scorer; test membership/distribution; environment/tools/data; aggregation/reporting rules; run logs/configuration		Code deliberate restrictions that affect evaluation or interpretation.	Items/tasks — evaluation prompts, cases, questions, or instances; labels/reference answers — ground truth or expected answers; rubric/judge/scorer — criteria or system used to judge outputs; test membership/distribution — which items are in the test set or how they are sampled; environment/tools/data — hidden execution state, tools, files, or test data; aggregation/reporting rules — how component results become the reported score; run logs/configuration — records or settings from evaluation runs; other — protected component not covered above.
Multi-select	Descriptive	**Restricted audience**	Who lacks ordinary access to the protected component?	Evaluated provider/team; benchmark participant/user; independent researcher; general public; customer/client		Select roles that are explicitly restricted.	Evaluated provider/team — organization or team whose model/system is being scored; benchmark participant/user — ordinary person or team using/submitting to the benchmark; independent researcher — outside researcher without evaluator privileges; general public — anyone without special access; customer/client — paying or contracted benchmark user; other — restricted audience not covered above; unclear — restriction is evident but the audience cannot be identified.
Single-select	Descriptive	**Secrecy lifecycle**	When, if ever, is the protected material intended to become accessible?	Indefinite/no planned release; rolling release/replacement; fixed or event-bounded release; until benchmark retirement; mixed by component; not reported		Use the stated policy. If components differ materially, select mixed.	Indefinite/no planned release — no release endpoint is stated; rolling release/replacement — old protected material is released or replaced as new material becomes private; fixed or event-bounded release — release follows a stated date, delay, competition, or evaluation round; until benchmark retirement — material is released only when that version/benchmark is retired; mixed by component — different protected components follow different lifecycles; not reported — no lifecycle policy located.
Multi-select	Descriptive	**Evaluation access**	How is evaluation performed while the protected material remains withheld?	Prediction submission to server; hosted task/evaluation API; evaluator runs the model/system; submitted code/container; secure/controlled environment; third-party platform/custodian		Code the operational evaluation mechanism.	Prediction submission to server — participant sends outputs for hidden scoring; hosted task/evaluation API — protected tasks are served through an API or hosted interface; evaluator runs the model/system — evaluator operates or queries the evaluated system; submitted code/container — participant submits executable code/container for organizer-side evaluation; secure/controlled environment — evaluation occurs in a restricted room, enclave, or similar setting; third-party platform/custodian — an external platform or custodian holds/runs protected material; other — access mechanism not covered above.
Multi-select	Descriptive	**Durability controls**	What mechanisms are used to preserve benchmark usefulness across repeated evaluation?	Submission/rate limits; refresh/rotation; fresh/on-demand generation; contamination/overlap monitoring; access controls; retirement/replacement; none documented		Code documented controls against exposure, adaptation, or repeated-use degradation.	Submission/rate limits — limits repeated attempts; refresh/rotation — periodically replaces protected items; fresh/on-demand generation — creates new evaluation material for runs/releases; contamination/overlap monitoring — checks for leakage, memorization, or dataset overlap; access controls— restricts who accesses protected material; retirement/replacement — withdraws degraded versions or sets; none documented — no durability control located.
				retirement	More "retirement" ---- ....---- more/less "durability"		
				access controls	More "access control" usually means less leakage, and less "leakage" usually means more "durability"		
		2. Validity evidence					
Multi-select	Comparison	**Claim framing**	Which elements needed to interpret the benchmark claim are explicitly stated?	Measured capability/construct; intended use or decision purpose; target domain/population; limits or out-of-scope interpretation; none documented		Code each element separately; do not use an overall Yes/Partial judgment.	Measured capability/construct — states what ability/property the score is meant to measure; intended use or decision purpose — states how the result is meant to be used; target domain/population — states the tasks, users, systems, or setting to which the result applies; limits or out-of-scope interpretation — states what the score should not be taken to mean; none documented — none of these elements is stated.
Multi-select	Comparison	**Task transparency**	What lets an outsider understand the task without seeing the live private set?	Representative public items/dev set; detailed task blueprint/schema; none documented		Examples show task form; a blueprint describes task structure without exposing live items.	Representative public items/dev set — public examples show the same basic task form as the private evaluation; detailed task blueprint/schema — documentation specifies task structure, inputs, outputs, or construction without revealing live items; none documented — neither public examples nor a sufficiently detailed blueprint is located.
Multi-select	Comparison	**Hidden-set coverage evidence**	What evidence supports assessing whether the hidden set covers the stated task space?	Sampling/source method documented; public-private comparability evidence; documented expert or independent review of hidden set; none documented		Code coverage/representativeness evidence, not merely task examples.	Sampling/source method documented — explains where hidden items come from or how they are sampled; public-private comparability evidence — directly compares public and private splits/distributions/difficulty; documented expert or independent review of hidden set — reviewers assess hidden-set coverage or appropriateness; none documented — no coverage evidence located.
Multi-select	Comparison	**Scoring transparency**	What lets an outsider understand how model outputs become reported scores?	Metric/formula public; rubric/judge criteria public; judge/scorer details public; aggregation/threshold rules public; controlled inspection only; none documented		Code scoring mechanics, not whether the score is substantively valid.	Metric/formula public — score computation is publicly specified; rubric/judge criteria public — qualitative grading criteria are public; judge/scorer details public — relevant judge model, scorer, prompt, or implementation details are public; aggregation/threshold rules public — combination, cutoff, weighting, or tie rules are public; controlled inspection only — scoring details can be inspected only under restricted access; none documented — no meaningful scoring details located.
		3. Integrity evidence					
Multi-select	Comparison	**Protocol rules**	Which consequential evaluation rules are documented?	Model/system eligibility or configuration; task selection/randomization; submission/retry/rerun rules; failure/exclusion handling; result reporting/publication rules; none documented		Code rules that could materially change a reported result.	Model/system eligibility or configuration — rules constrain what system/configuration may be evaluated; task selection/randomization — rules specify how tasks are selected, ordered, or randomized; submission/retry/rerun rules — rules govern repeated attempts or reruns; failure/exclusion handling — rules specify how invalid, failed, or excluded cases are treated; result reporting/publication rules — rules govern which results are reported or published; none documented — none of these procedural rules is located.
Multi-select	Comparison	**Ex-ante fixation**	What evidence indicates consequential rules or content were fixed before the relevant results?	Dated/versioned protocol; preregistration or immutable commitment; third-party/escrow confirmation; none documented		A protocol description alone is not ex-ante evidence unless it is dated/versioned.	Dated/versioned protocol — relevant rules/content are tied to a public date or version before the result; preregistration or immutable commitment — preregistration, hash, signed commitment, or similar prevents silent later changes; third-party/escrow confirmation — an external party or escrow held/confirmed the relevant material before results; none documented — no evidence of pre-result fixation located.
Multi-select	Comparison	**Run provenance**	Which details identify the evaluation run that produced a reported score?	Benchmark version/snapshot; model version/checkpoint; inference/system configuration; evaluator/scorer version; run date; task subset/seed/session; none documented		Code only details that can be tied to the reported score.	Benchmark version/snapshot — exact benchmark release or snapshot is identifiable; model version/checkpoint — exact evaluated model/version is identifiable; inference/system configuration — material evaluation settings are recorded; evaluator/scorer version — evaluation harness, judge, or scorer version is identifiable; run date — evaluation date/time is recorded; task subset/seed/session — evaluated subset, random seed, or run/session identifier is recorded; none documented — none of these run identifiers is tied to the score.
Multi-select	Comparison	**Run audit evidence**	What retained evidence could be used to check whether the reported run followed the stated procedure?	Public outputs/logs/run record; controlled-access run records; evaluator or third-party replay possible; none documented		This is evidence available for checking a run, not proof that an audit occurred.	Public outputs/logs/run record — outsiders can inspect outputs, logs, or a run record; controlled-access run records — relevant run records exist but require restricted access; evaluator or third-party replay possible — documented process allows the run/result to be replayed by the evaluator or another party; none documented — no retained evidence for checking the run is located.
		4. Neutrality & access					
Single-select	Comparison	**Participation-rule visibility**	Who can inspect the rules governing eligibility and access to the benchmark?	Public; participants/customers or on request only; no rules located		Code visibility of the rules, not whether they are fair.	Public — eligibility/access rules are openly available; participants/customers or on request only — rules are available only after joining, paying, applying, or requesting access; no rules located — no eligibility/access rules were found.
Single-select	Comparison	**Differential-treatment disclosure**	Does the documentation state whether participants/providers may receive different access or evaluation treatment?	Uniform treatment explicitly stated; differential/privileged treatment disclosed or permitted; not reported		Code the stated policy only; do not infer equal treatment from silence.	Uniform treatment explicitly stated — documentation says comparable participants receive the same relevant access/conditions; differential/privileged treatment disclosed or permitted — documentation states or allows materially different treatment; not addressed — documentation does not state whether treatment is uniform or differential.
Multi-select	Comparison	**Evaluator–provider ties**	What material relationships between the evaluator/operator and evaluated providers are documented?	Same organization/affiliate; funding/sponsorship; customer/vendor relationship; research/technical partnership; privileged model/data/API access; other material tie; explicit no material ties; no relationship information located		Positive tie codes are mutually exclusive with “explicit no material ties” and “no relationship inform	Same organization/affiliate — evaluator and evaluated provider share ownership or organizational control; funding/sponsorship — evaluated provider funds or sponsors the evaluator/benchmark; customer/vendor relationship — one party buys services/products from the other; research/technical partnership — documented collaboration relevant to the benchmark or evaluated system; privileged model/data/API access — evaluator receives access not generally available to comparable parties; other material tie — relevant tie not covered above; explicit no material ties — documentation explicitly states none; no relationship information located — no relevant disclosure was found.
		5. Verification regime					
Single-select	Comparison	**External verification status**	What evidence shows that a party outside the benchmark operator and evaluated provider actually performed a relevant c	Verifier-authored or other external evidence located; operator reports external verification only; no external verification located		An available audit mechanism does not count unless an external check actually occurred.	Verifier-authored or other external evidence located — an outside verifier or independent source documents that a check occurred; operator reports external verification only — only the benchmark operator states that an outside check occurred; no external verification located — no evidence of an external check was found.
Conditional singl	Comparison	**Verifier relationship**	If external verification occurred, how is the verifier related to the evaluator/evaluated parties?	Independent external; external with a material tie; relationship unclear; N/A		Base the classification on documented organizational, financial, or governance relationships.	Independent external — no material organizational, financial, or governance tie is documented; external with a material tie — verifier is outside the operator but has a relevant documented tie; relationship unclear — external verification occurred but independence cannot be determined; N/A — no external verification occurred/was located.
Conditional multi	Comparison	**Verification scope**	If external verification occurred, what did the verifier actually inspect or test?	Protected content/coverage; scoring/rubric; run/execution records; access/participant records; governance/conflicts; public documentation only; N/A		Code actual inspection scope, not hypothetical access.	Protected content/coverage — hidden items, sampling, or representativeness; scoring/rubric — metric, rubric, judge, or scoring process; run/execution records — logs, configurations, outputs, or execution procedure; access/participant records — participation, access, or treatment records; governance/conflicts — funding, relationships, or governance arrangements; public documentation only — verifier reviewed only material already public; other — verified scope not covered above; N/A — no external verification occurred/was located.
Conditional singl	Comparison	**Findings disclosure**	If external verification occurred, what findings are available to benchmark users?	Detailed public findings; public summary/conclusion; verification existence only; restricted/confidential findings; no findings located; N/A		Distinguish disclosure of findings from disclosure that verification occurred.	Detailed public findings — substantive methods and findings are publicly available; public summary/conclusion — only a high-level result or conclusion is public; verification existence only — public information says verification occurred but gives no finding; restricted/confidential findings — findings exist but are not public; no findings located — verification occurred but no findings disclosure was found; N/A — no external verification occurred/was located.